Ziad Abou Jaoude

MLOps / Platform Engineer

ziad.aboujaoude@example.com | 76 119 508 | linkedin.com/in/ziadaj

SUMMARY

Platform engineer who has spent five years making other people's models run. I own serving, scaling and observability for ML workloads. I am not the person who trains the model.

EXPERIENCE

MLOps Engineer, Argan AI - Beirut and Remote

11/2021 - Present

- Own the model serving platform: 34 models in production on Kubernetes behind Triton Inference Server and KServe, roughly 900 requests per second at peak.
- Cut GPU spend 44 percent by introducing dynamic batching and moving eight models to int8 via TensorRT, holding p99 within the 250 millisecond SLO.
- Built the deployment pipeline that takes a model from an MLflow registry entry to a canary rollout with automatic rollback on latency or error-rate regression.
- Run the feature and embedding infrastructure, including a 30 million vector index in Milvus and the reindexing job that rebuilds it without downtime.
- Wrote the model monitoring stack: input drift detection, prediction distribution alerts, and per-model latency dashboards in Grafana.
- Set up the GPU node pools, autoscaling and spot instance handling on EKS.

DevOps Engineer, Halcyon Systems - Beirut

07/2019 - 10/2021

- Kubernetes, Terraform and CI/CD for twelve microservices.
- Built the observability stack: Prometheus, Grafana, Loki, and on-call runbooks.

EDUCATION

BE Computer and Communications Engineering, American University of Beirut, 2019

SKILLS

Python, Go, Kubernetes, Docker, Terraform, KServe, Triton, TensorRT, ONNX, MLflow, Milvus, pgvector, Prometheus, Grafana, AWS, EKS.

HONEST SCOPE

I write Python for tooling and pipelines, not for modelling. I have run PyTorch training jobs on the cluster and debugged them when they OOM, but I have never designed a training run or chosen a loss function. I have never built an offline evaluation set; I consume the metrics the DS team produces. Transformers I understand operationally, not mathematically.